

Mandate, full documentation

Team 10, Raingentic Commerce Hackathon NYC, 2026-08-08/09.

1. What it is

An agent gets a **scoped virtual Rain card** bound to the exact purchase order it negotiated: vendor, SKU, price, quantity, expiry. Deterministic code checks the declared order against the real record before the card is ever created. Any mismatch means no card, not a decline, an instrument that never comes into existence.

Rain’s own products bound *how much* an agent spends and *where*. Mandate binds *why*, at the same moment, one step earlier.

2. The architecture

1 TASK	agent is given a job	
1b QUOTE	buyer gets a vendor quote (competing quotes if time allows)	
	v	
2 PROPOSE	agent declares the accepted PO	
		{ poNumber, vendor, sku, unitPrice, quantity, quoteExpiry }
	v	
3 VERIFY	deterministic checks against the record	<-- no model here
		rules are versioned CONFIG, not code
	pass fail	
		\
		\--> REFUSE no card. plain-English reason. logged.
	v	
4 ISSUE	Rain issues a scoped virtual card, bound to exactly this PO	
	v	
5 SETTLE	the purchase happens on that card	
	v	
6 RECORD	PO + checks + rule version + card + outcome, append-only	
7 REVOKe	card deactivated once the job is done (if the endpoint supports it)	

Monad: each **rule version** is hashed and anchored on testnet. This is what proves the rules were not edited after the fact to fit a history someone already had, closing the one real hole in an otherwise-airtight audit story.

Six deterministic checks decide stage 3: the PO exists and is accepted, it’s still open, the amount matches the quote, the merchant matches the quote, it’s within budget, and no card has already been issued for this exact PO (idempotency, the check that carries the demo).

3. Why this company, specifically

Rain is not a generic payments API. It is a **Visa and Mastercard Principal Member**, meaning it issues cards directly rather than reselling someone else's licence, live across 175 million merchant locations in 220+ countries, backing 100+ programs on \$250M raised. Its newest product, shipped this June, is the **Agent Control Layer**: programmatic guardrails, amount limits, merchant allowlists, spend intervals, enforced *"at card issuance and transfer initiation rather than applied after the fact."*

That last phrase is Mandate's whole thesis. Rain already committed, publicly, to enforcement-before-the-fact as the right design. Mandate is the same principle carried one layer higher, from "can this agent spend this much, here" to "is this specific declared reason for spending actually true." It does not compete with the Agent Control Layer. It sits on it.

Concretely, the build uses:

- **Rain's collateral-backed card issuance** (`issueScopedCard` in `lib/rain-client.ts`), the same primitive shown live in Rain's own "Wednesday We Wear Pink" reference demo during the workshop: onramp → store stablecoins → user authorizes agent → **scoped virtual card issued, locked to merchant and amount** → agent completes the transaction → **card retired automatically once the job is done**. Mandate's issue/settle/revoke stages are that exact flow.
- **The team's own collateral contract**, pre-provisioned (`RAIN_COLLATERAL_CONTRACT_ID`), meaning the KYC and contract-deployment steps Rain normally requires were already done for Team 10 by Rain, and the build starts at the step that matters, issuance.

4. Why this hackathon, specifically

Three of the five judges work at Rain: Ross Basri (product lead, launched Rain Rewards), Farhan Khwaja (engineer, high-throughput transactional systems, custody infrastructure), Juan Blanco (data engineer, has personally won Ethereum hackathons in Amsterdam and Paris, so he will recognize a staged demo on sight). The other two are Siggy Bilstein (Engineering Manager at Cursor, works on an agent-native git forge) and Jarrod Watts (AI Engineering Lead at Monad, works specifically on agent orchestration).

This shaped four concrete build decisions:

- **The lead demo moment is unfalsifiable, not scripted.** Issue a card for a PO, then submit the identical PO again. It is refused by the idempotency check, reading the record the first run itself wrote. Nothing is authored for the demo. This exists because Juan Blanco would spot a hand-authored "bad agent" fixture in seconds.
- **No LLM anywhere in the verify path.** Every check in `lib/verify.ts` (B's side) is a pure function. This is what makes replay meaningful, editing a rule and re-running history only proves something if the thing being re-run is deterministic. Farhan builds transactional systems for a living; "trust me, the agent judged it fairly" does not survive contact with that judge.
- **The rule-version Monad anchor, not a per-decision one.** Anchoring every decision is a gesture. Anchoring each rule version is structural, it is what proves the rules were not quietly rewritten to fit a history after the fact, and it is the honest answer to Jarrod's own stated bounty bar: *"the chain has to matter, would it break at 15 second finality or 50 cents a transaction?"* You can only afford to anchor rule history because Monad is this cheap.

- **The negotiation stage stays deliberately thin**, one vendor comparison rather than a multi-agent haggling performance, specifically because Jarrod’s actual job is agent orchestration and a shallow simulated negotiation is the easiest thing in the room for him to see through.

5. How it covers all four submission categories

The event’s tracks are not mutually exclusive submissions, one build can qualify for several with different framing. Rain’s own workshop slide named four shapes for “Autonomous spend”: **Autonomous spend, Global money movement, Treasury and payouts, Agent negotiation**. Mandate maps to two of Rain’s own four, plus the general track and the Monad bounty:

Category	How Mandate qualifies
Best use of Rain	Card issuance is the enforcement point, matching Rain’s own architecture and their published “at issuance, not after the fact” principle exactly
General track, “agents actually move money”	Real settlement through a real scoped virtual card, not a simulation
Agent negotiation	The quote stage produces the accepted PO the card is bound to; described honestly as quote-selection rather than an inflated multi-round negotiation
Monad bounty	Real testnet transactions (rule-version hashes), where the chain’s low cost is structurally why the audit claim can hold at all, not a decorative add-on

One pipeline, one architecture, two owners. Not four separate projects chasing four prizes.

6. What was deliberately not built, and why

Documented in full in [IDEAS.md](#) , [THE-IDEA.md](#) , and [ABI-LESSON.md](#) . In short:

- **Delegated budget trees** — too close to Rain’s own program-level caps, would read as rebuilding their product rather than extending it.
- **Agent Underwriting** (a second collateral pool taking on liability) — needs Rain identities the team was not issued. Kept as one forward-looking sentence in the close.
- **Burn Rate streaming limits** — needs an unconfirmed Rain capability (live limit updates post-issuance) plus a Monad streaming contract from scratch. Two coupled unknowns on one afternoon.
- **Four-Eyes consensus voting** — rejected on principle, not just time. LLM agents voting reintroduces a model into the verify path, which directly undoes the reason replay works at all.
- **A per-decision Monad fee, a live budget meter, a judge-editable form** — all real, all kept as optional additions, none load-bearing, all cut before anything in the core loop.

This scoping follows directly from a prior hackathon loss (Pulse Foundry × ABI Frameworks, 2026-06-28): the team built roughly 85% of the winning system’s substance and still lost, to a team whose actual edge was a

versioned, biller-editable rules config and a “re-run instantly” toggle. Mandate’s replay feature is that same idea, built properly, from the start, rather than bolted on after losing to it once.

7. Stack and repo layout

Next.js 14 App Router, TypeScript, Tailwind. Deployed on Vercel, required, locally-hosted submissions are disqualified per the event’s own rules.

```
lib/
  types.ts      shared PurchaseOrder / VendorQuote / CheckResult shapes (A + B contract)
  money.ts      integer cents, never floats
  rain-client.ts Rain API wrapper (A)
  quotes.ts     vendor quote fixtures + selection (A)
  agent.ts      task -> PurchaseOrder (A)
  verify.ts     the six checks, rule versioning, replay (B)
  db.ts         append-only decision log storage (B)
app/
  api/rain/ping/ the first-authenticated-call milestone
  api/purchase/  the 17:00 join: agent proposes -> verify -> issue
  api/replay/    re-run history against an edited rule version (B)
docs/           this file, and the full planning trail
```

Repo: `github.com/theomthakur/raingentic-team10`, public per the event’s rules, `.env.local` never committed, only `.env.local.example`.